

# The Process and Standard File Descriptors (fds)

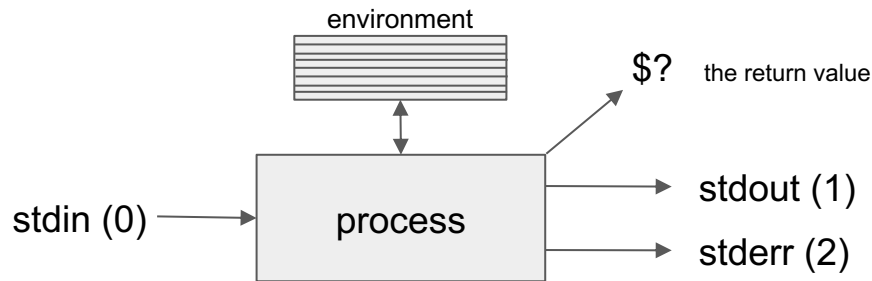

- A process is a running program
- From the CLI:

- `java program arg0 arg1 arg2`
- `program arg0 arg1 arg2`

- From within the program

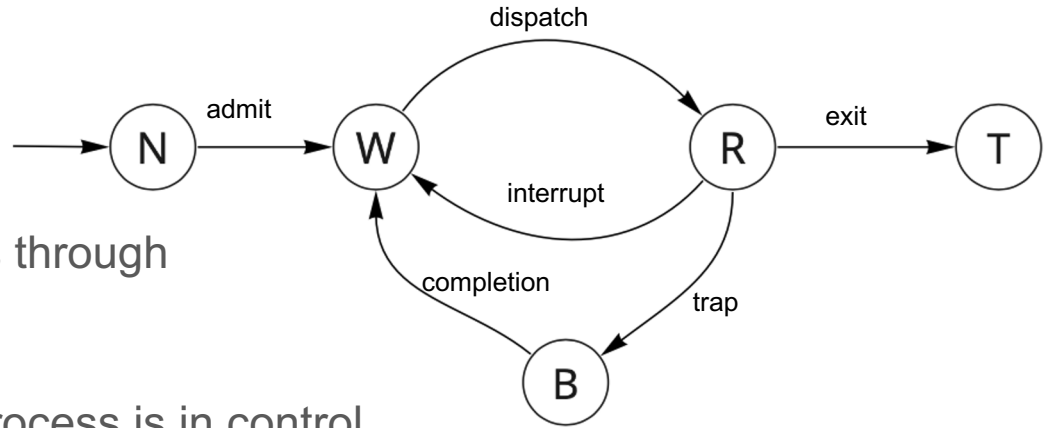
- `public static void main(String[] args)` //Java
- `void main(int argc, char* argv[], char** env)` // C

- Variables within memory:



Java Parlance:  
`System.in` == `stdin`  
`System.out` == `stdout`  
`System.err` == `stderr`

# Process Status Diagram



- Control of the computer moves through a well-defined cycle
- At any point in time, a single process is in control
  - loosely speaking a process is equivalent to a program
- Transitions:
  - admit: A request is made to allow your program to content for control
  - dispatch: Your program is given control
  - exit: Your program asserts that it is done
  - interrupt: The OS seizes control
  - trap: Your program (implicitly or explicitly) requests a service to be performed
  - completion: The request is satisfied

# MIPS ISA Architecture: OS interface


```
# Print the integer '1'
# Macro: print_di %imm
li $a0, %imm
li $v0, 1
syscall
```

- Service Requests to the operating system via the 'syscall' instruction

Service Name	\$v0	input: \$a0..\$a3	output: \$v0..\$v1
print integer	1	\$a0 = value	none
read integer	5	none	\$v0 = value
malloc	9	\$a0 = size	\$v0 = buffer address
exit	10	none	none
file read	14	\$a0 = fd, \$a1 = buffer address \$a2 = num bytes	\$v0 = bytes read -1 == error 0 == eof

# Execution of your program

1. Invoke the program:
2. Wait to use the CPU
3. Execute for as long as you can -- Until
  - (Exit) You are done
  - (Interrupt) You get interrupted by some outside force
  - (Trap) You need help because you made an error or you requested it
4. If you were interrupted, goto Step 2
5. If you trap, and then goto Step 2
  - recover from the error, or
  - obtain the requested server



Analogy: Driving your Car from LA to Vegas

# Interrupts and Traps: (results in the kernel seizing control)

- Interrupts are asynchronous events
    - such events occur outside of your process/program
    - such events may or may not be associated with your program
    - Examples:
      - data has arrived on the NIC
      - a disk request for a different process has been completed
  - Traps are synchronous events
    - such events occur inside of your process/program
    - some events are error conditions, e.g.,
      - division by zero
      - invalid or illegal memory access
    - some events are requests, e.g.,
      - read/write from a file
      - create a child process
  - Exits are a specific type of trap that results in a different flow through the PSD
- For speed, traps are to be avoided!

# "read" system call

- You need to allocate a buffer, a block of memory.

```
byte buffer[8];  
int * p = &buffer;
```

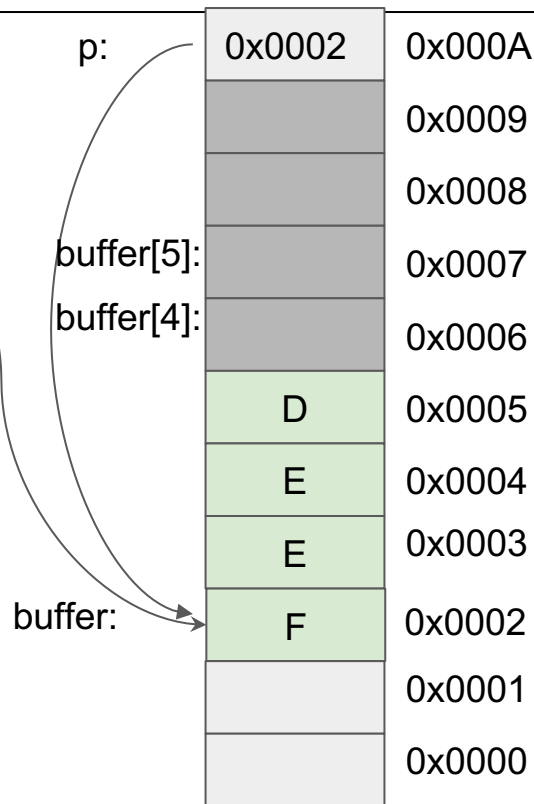
- Make a read request to the OS, providing:
  - the identifier of the file to read
  - the location of the buffer
  - the number of bytes to read

```
retval = read(fd, &buffer, 8);
```

- What are the values passed to read?
- Value of retval informs what happened.
  - retval == -1: error
  - retval == 0: end of file
  - retval <= 8: number of bytes read
- Cast the code
  - retval = read(fd, (void \*) &buffer, 8);

```
int retval;  
int fd;  
fd = open("/home/steve/filename", O_RDONLY);
```

p: 0x0002



```
Java: File fd = new Scanner("/home/steve/filename");  
buffer = fd.nextByte(); // in C: read(fd, &buffer, 1)
```

# Reading a block of 10 bytes!

- Java Example:

```
byte header[10];
stdin = new Scanner(System.in);
for( i = 0; i < 10 ; i++ ) {
    header[i] = stdin.nextByte();
}
```

- The Scanner class only handles primitive types
- We are not in a position to reimplement it.
- The OS knows nothing about my Java class
- Consequently, this results in 10 systems calls

- Equivalent C Example

```
byte header[10];
for( i = 0; i < 10 ; i++ ) {
    header[i] = (byte) getchar();
}
```



The diagram consists of a horizontal arrow pointing from the C code block to the read call. The arrow originates from the right side of the C code block and points to the left side of the read call box.

```
retval = read(fd, &header, 10);
```

# A More Efficient Approach

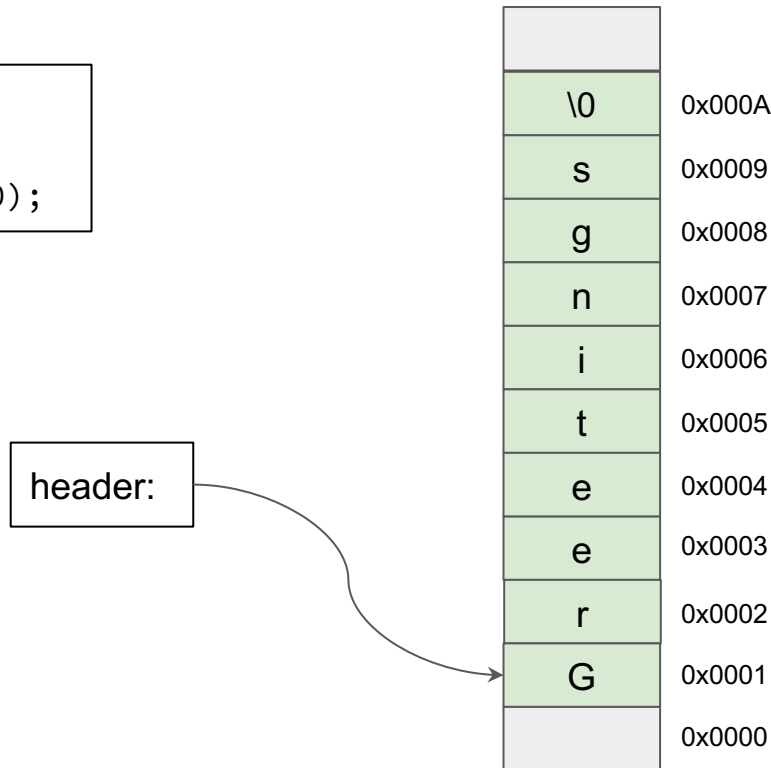
- Via 'read' system call (1 trap):

```
byte header[10];  
  
retval = read(STDIN_FILENO, (void *) &header, 10);
```

- Java Example (10 traps):

```
byte header[10];  
stdin = new Scanner(System.in);  
for( i = 0; i < 10 ; i++ ) {  
    header[i] = stdin.nextByte();  
}
```

- Read does not care what it is reading
- This results in 1 system call
- But we need to understand pointers: \* and &
- Moreover we need to cast our variables






# Variables: Names and Values

- We all know what variables are, right!
- A variable has two values: a lval and a rval
- Consider the following assignment:

`c = c + 1`



1. determine the lval for “c”
2. extract the rval of “c” at MEM[lval]
3. add 1 to that value
4. Write the value from #3 into MEM[lval]

- The lval of `c` is the address in memory
- The rval of `c` is the value located at that address in memory
- The rval can be a "integer", "float", "char", etc.
- If the rval is an address, than it is known as a "pointer"

`int * p = &c;`

`p:`

?

- Hold that thought! `(* p) ==`

	0x000A	0
	0x0009	0
d:	0x0008	5
	0x0007	0
c:	0x0006	2
	0x0005	
b:	0x0004	6
p:	0x0003	?
	0x0002	0
a:	0x0001	1
	0x0000	42

# Memory

- We all know what an array is right!
  - Memory is just an array of integers (from 0..255):
    - `mem[ index ] = value`
- Do you know what an associative array is?
  - It's just an array that stores both the lval and rval of a variable:
    - `array[ "name" ] = value;   mem[ "steven" ] = 32`
  - You use "name" to lookup the appropriate index
- Consider the memory to the right

		0x8000 000A
	0	0x8000 0009
sasank:	3	0x8000 0008
	0	0x8000 0007
shant:	37	0x8000 0006
		0x8000 0005
steven:	32	0x8000 0004
syndey:	?	0x8000 0003
	45	0x8000 0002
tyler:	1	0x8000 0001
	0	0x8000 0000

\*\*\* p == 42

# Variables: Names and Values


- We all know what variables are, right!
- A variable has two values: a lval and a rval
- Consider the following assignment:

`c = c + 1;`

- The lval of c is the address in memory
- The rval of c is the value located at that address in memory
- The rval can be a "integer", "float", "char", etc.
- If the rval is an address, then it is known as a "pointer"

`int * p = &c;`

- Hold that thought! `(* p) == 2`



	0x000A	0
	0x0009	0
d:	0x0008	5
	0x0007	0
c:	0x0006	2
	0x0005	
b:	0x0004	6
p:	0x0003	?
	0x0002	0
a:	0x0001	1
	0x0000	42

- What is the address c? What is the value of p? What is the address of p?